

In [4]:

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
from sklearn.metrics import accuracy_score
import seaborn as sns
from sklearn.metrics import confusion_matrix

#importing dataset
df = pd.read_csv('/content/Global_PowerPlant_Database.csv')
df
```

Out[4]:

	country_id	country_name	powerplant_name	gppd_idnr	capacity_mw	latitude	longitude
0	1	Australia	Albany Wind Farm	AUS0000065	216.0	-350641	1177977
1	1	Australia	Appin (Mine)	AUS0000264	556.0	-342082	1507722
2	1	Australia	Awaba	AUS0000049	11.0	-330250	1515497
3	1	Australia	Bairnsdale	AUS0000081	940.0	-378429	1475639
4	1	Australia	Ballarat Base Hospital	AUS0000113	30.0	-375596	1438469
...
6812	3	United States of America	Zion Energy Center	USA0055392	5967.0	424776	-878950
6813	3	United States of America	Zion Landfill Gas to Energy Facility	USA0056871	70.0	424803	-878861
6814	3	United States of America	Zorn	USA0001368	180.0	382803	-857023
6815	3	United States of America	Zotos International WPGF	USA0057648	34.0	428869	-769683
6816	3	United States of America	eBay - South Jordan	USA0059281	98.0	405606	-1120478

6817 rows x 17 columns

cleaning data

In [49]:

```
print(df.shape)
```

(6817, 17)

```
In [50]:  
  
#Dropping the missing rows.  
df_dropped = df.dropna(how = 'any')  
#there is no null value in this dataset
```

```
In [6]:  
  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 6817 entries, 0 to 6816  
Data columns (total 17 columns):  
#   Column                Non-Null Count  Dtype  
---  ---  
0   country_id            6817 non-null   int64  
1   country_name          6817 non-null   object  
2   powerplant_name       6817 non-null   object  
3   gppd_idnr             6817 non-null   object  
4   capacity_mw           6817 non-null   float64  
5   latitude              6817 non-null   int64  
6   longitude             6817 non-null   int64  
7   primary_fuel_id       6817 non-null   int64  
8   primary_fuel          6817 non-null   object  
9   source                6817 non-null   object  
10  url                   6817 non-null   object  
11  geolocation_source    6817 non-null   object  
12  generation_gwh_2013   6817 non-null   float64  
13  generation_gwh_2014   6817 non-null   float64  
14  generation_gwh_2015   6817 non-null   float64  
15  generation_gwh_2016   6817 non-null   float64  
16  generation_gwh_2017   6817 non-null   float64  
dtypes: float64(6), int64(4), object(7)  
memory usage: 905.5+ KB
```

```
In [7]:  
  
df.describe()
```

Out[7]:

	country_id	capacity_mw	latitude	longitude	primary_fuel_id	generation_gwh
count	6817.000000	6.817000e+03	6817.000000	6.817000e+03	6817.000000	6.817000e+03
mean	2.886460	1.784192e+15	363351.969781	-7.945646e+05	4.396362	4.176000e+03
std	0.399552	7.520842e+15	140160.924371	5.811068e+05	3.303483	1.135470e+03
min	1.000000	1.000000e+01	-428883.000000	-1.717124e+06	1.000000	-6.000000e+03
25%	3.000000	5.000000e+01	341452.000000	-1.105739e+06	2.000000	2.661000e+03
50%	3.000000	3.640000e+02	393740.000000	-9.017970e+05	4.000000	5.715800e+03
75%	3.000000	2.547000e+03	426260.000000	-7.711890e+05	6.000000	1.954700e+03
max	3.000000	7.900000e+16	712920.000000	1.531236e+06	14.000000	7.919980e+03



In [8]:

```
df1 = df.append(df.iloc[1:6817,:])  
  
df1.duplicated()
```

Out[8]:

```
0      False  
1      False  
2      False  
3      False  
4      False  
...  
6812    True  
6813    True  
6814    True  
6815    True  
6816    True  
Length: 13633, dtype: bool
```

training and testing data

In [9]:

```
# x and y values for test-train datas.  
y = df.primary_fuel_id.values  
x = df[['capacity_mw', 'latitude', 'longitude', 'country_id', 'generation_gwh_2013', 'gene
```

In [10]:

```
y
```

Out[10]:

```
array([1, 2, 3, ..., 2, 1, 2])
```

```
In [11]:
x
```

Out[11]:

	capacity_mw	latitude	longitude	country_id	generation_gwh_2013	generation_gwh_2014	g
0	216.0	-350641	1177977	1	8.959530e+15	1.026430e+16	
1	556.0	-342082	1507722	1	2.048040e+16	2.382940e+16	
2	11.0	-330250	1515497	1	7.655280e+15	8.338610e+15	
3	940.0	-378429	1475639	1	1.324570e+16	1.548860e+16	
4	30.0	-375596	1438469	1	4.194440e+15	4.213060e+15	
...
6812	5967.0	424776	-878950	3	0.000000e+00	6.365600e+16	
6813	70.0	424803	-878861	3	4.359200e+04	4.913400e+04	
6814	180.0	382803	-857023	3	2.030000e+02	7.800000e+01	
6815	34.0	428869	-769683	3	3.671000e+16	2.460150e+05	
6816	98.0	405606	-1120478	3	1.296800e+04	5.127600e+04	

6817 rows × 9 columns



In [12]:

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.2,random_state =
#test_size=0.2 means %20 test datas, %80 train datas
method_names = []
method_scores = []
print(x_train)
print(y_train)
```

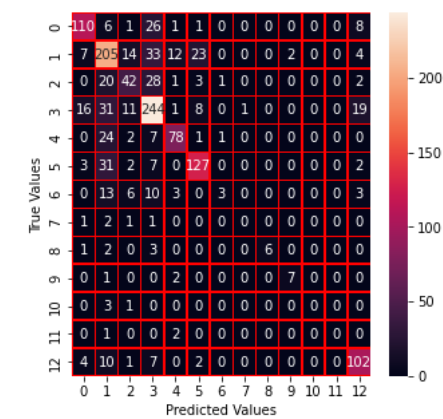
	capacity_mw	latitude	longitude	country_id	generation_gwh_201
3 \					
6443	1.520000e+16	391992	-963086	3	0.000000e+0
0					
1620	4.000000e+01	335683	-1149181	3	5.403000e+1
6					
2820	1.720000e+02	421212	-1225479	3	5.696000e+0
3					
1643	1.375600e+04	360875	-1150507	3	6.544230e+0
5					
2877	1.656000e+03	327683	-994228	3	4.821470e+0
5					
...	
...					
3772	1.100000e+01	407576	-1214145	3	4.893000e+0
3					
5191	2.490000e+02	262697	-978670	3	3.377000e+1
6					
5226	9.120000e+03	450038	-747994	3	6.672860e+1
5					
5390	2.000000e+01	372875	-800771	3	1.140000e+0
2					
860	2.020000e+02	358950	-1195108	3	4.216100e+0
4					
	generation_gwh_2014	generation_gwh_2015	generation_gwh_2016	\	
6443	-3.280000e+02	-7.330000e+15	-7.460000e+02		
1620	5.543000e+03	5.376000e+03	9.016000e+03		
2820	6.442300e+04	5.448100e+04	5.225500e+04		
1643	5.184870e+15	8.991210e+15	9.075430e+15		
2877	5.144550e+05	4.413090e+05	4.594550e+05		
...		
3772	4.772000e+03	4.284000e+03	4.032000e+03		
5191	2.342000e+16	4.210000e+02	2.192000e+03		
5226	7.052696e+06	6.932867e+06	7.088396e+06		
5390	1.020000e+16	1.040000e+16	1.710000e+02		
860	4.202500e+04	4.131400e+04	4.125500e+04		
	generation_gwh_2017				
6443	-6.760000e+15				
1620	8.962000e+03				
2820	4.368500e+04				
1643	8.439040e+05				
2877	4.366440e+05				
...	...				
3772	5.393000e+03				
5191	1.937000e+03				
5226	7.632069e+06				
5390	9.800000e+01				
860	3.806300e+04				

[5453 rows x 9 columns]
[2 13 4 ... 4 6 13]

Random Forest Classification


```
In [13]:  
  
from sklearn.ensemble import RandomForestClassifier  
RF = RandomForestClassifier()  
RF.fit(x_train,y_train)  
print("random forest classification test Accuracy: ",RF.score(x_test,y_test))  
method_names.append("Random Forest Classification")  
method_scores.append(RF.score(x_test,y_test))  
  
#Confusion Matrix  
y_pred = RF.predict(x_test)  
conf_mat = confusion_matrix(y_test,y_pred)  
#Visualization Confusion Matrix  
f, ax = plt.subplots(figsize=(5,5))  
sns.heatmap(conf_mat,annot=True,linewidths=0.5,linecolor="red",fmt=".0f",ax=ax)  
plt.xlabel("Predicted Values")  
plt.ylabel("True Values")  
plt.show()
```


random forest classification test Accuracy: 0.6774193548387096



Decision Tree Classification

```
In [14]:  
  
from sklearn.tree import DecisionTreeClassifier  
dec_tree = DecisionTreeClassifier()  
dec_tree.fit(x_train,y_train)  
print("Decision Tree Classification test Accuracy: ",dec_tree.score(x_test,y_test))  
method_names.append("Decision Tree")  
method_scores.append(dec_tree.score(x_test,y_test))  
  
#Confusion Matrix  
y_pred = dec_tree.predict(x_test)  
conf_mat = confusion_matrix(y_test,y_pred)  
#Visualization Confusion Matrix  
f, ax = plt.subplots(figsize=(5,5))  
sns.heatmap(conf_mat,annot=True,linewidths=0.5,linecolor="red",fmt=".0f",ax=ax)  
plt.xlabel("Predicted Values")  
plt.ylabel("True Values")  
plt.show()
```

Decision Tree Classification test Accuracy: 0.5447214076246334



hypothesis testing

In [42]:

```
from scipy.stats import zscore, norm
#Calculating the z-score for the dataset
z_score = zscore(df.capacity_mw)

#Calculating the p-value
p_value = norm.sf(abs(z_score))

#Printing the results
print ("The z-score is :", z_score)
print ("The p-value is :", p_value)
```

```
The z-score is : 0      -0.23725
1      -0.23725
2      -0.23725
3      -0.23725
4      -0.23725
...
6812   -0.23725
6813   -0.23725
6814   -0.23725
6815   -0.23725
6816   -0.23725
Name: capacity_mw, Length: 6817, dtype: float64
The p-value is : [0.40623127 0.40623127 0.40623127 ... 0.40623127 0.406
23127 0.40623127]
```

data visualization

In [30]:

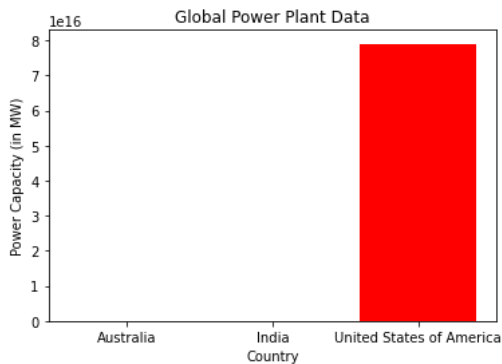
```
import matplotlib.pyplot as plt

# Input Data
country = df.country_name
power_capacity = df.capacity_mw

# Plotting a bar graph
plt.bar(country, power_capacity, color='red')

# Naming the x-axis
plt.xlabel('Country')
# Naming the y-axis
plt.ylabel('Power Capacity (in MW)')
# Graph Title
plt.title('Global Power Plant Data')

# Show the plot
plt.show()
```



In [36]:


```
import matplotlib.pyplot as plt

# Input Data
powerplant = df.powerplant_name
fuel = df.primary_fuel

# Plotting a bar graph
plt.bar(powerplant, fuel, color='blue')

# Naming the x-axis
plt.xlabel('Power Plant')
# Naming the y-axis
plt.ylabel('Fuel type')
# Graph Title
plt.title('Global Power Plant Data')

# Show the plot
plt.show()
```



In [48]:

```
import matplotlib.pyplot as plt

# Input Data
powerplant = df.primary_fuel_id
fuel = df.capacity_mw
# Plotting a bar graph
plt.scatter(powerplant, fuel, color='blue')

# Naming the x-axis
plt.xlabel('primary_fuel_id')
# Naming the y-axis
plt.ylabel('capacity')
# Graph Title
plt.title('Global Power Plant Data')

# Show the plot
plt.show()
```

